

Moral Deliberation and the Prospect of Parity

Derek Shiller

Abstract: If our actions have widespread and varied effects on the far future, then it is likely that they will make the world better and worse in many different sorts of ways. It is often hard to make normative comparisons between options that excel in different ways, due to imprecision in tradeoffs between values. There is reason to think that because of this, the total consequences of the actions that we routinely deliberate between are generally on a par. This should push us to reevaluate our approach to making moral decisions. If our actions are likely to have no positive or negative effect on the future as a whole, we should give less weight to their predictable effects. Though consequences matter, the causal interconnectedness of events gives us reason to favor non-consequentialist reasons for acting.

Imagine you spy a lost wallet on a walk in your local park. Returning the wallet will mildly inconvenience you and disproportionately help the owner. Suppose that through a circuitous and unpredictable chain of events, returning the wallet will also cause a beloved family dog to get loose and wander alone and confused for three days before being found; a few rats to find their way into the basement of an apartment building, multiplying before all dying to a slow and painful poison; a couple to plan a romantic evening, leading to the unplanned conception of a person whose life complicates theirs; a young woman to settle on a career as a musician, which she ultimately finds much less fulfilling than her alternative

options; an amateur mathematician to prove an obscure theorem of set theory; a local politician to successfully bribe his way out of corruption charges; a poem written by a refugee to be disseminated to a wide audience, inspiring some high school students and boring others; and twenty-six square kilometers of the Amazon rain forest to be burned, causing the extinction of an undiscovered species of beetle; among many other things.

Suppose further that careful scrutiny of such a long list of diverse effects would leave you and other reasonable people ambivalent about whether the world would be made better or worse overall by the myriad effects of the return of the wallet. It is generally hard to balance so many different things, but suppose that the difficulty in this case comes not from the cognitive load, but because the values involved are too different for precise comparisons. Knowing this, how should you think about your obligations with the wallet?

It is conceivable that most of the choices that face us in moral deliberation are like this: the lie you may tell your spouse, the change you may give to the panhandler, the gas guzzler you may choose to drive. Even if the most obvious consequences favoring one alternative, the unpredictable effects may make it a wash. All things considered, it may turn out that when all its consequences¹ are counted, no option seems to leave the world any better than any other.

I'll offer an argument for that conclusion. If the conclusion is correct, it surely must have some bearing on how we think about our moral obligations. Consequences loom large in moral deliberation. If full information would leave us indifferent between the total consequences of our actions, we should probably not treat predictable consequences so seriously.

¹I interpret the *consequences* of an action to include all of those events that counterfactually depend on the action taking place – i.e. that wouldn't have occurred if the action hadn't. Some of the consequences of an action will involve long causal chains and depend upon the activities of other agents.

The plan for this paper is as follows. In the first section, I argue that when faced with a choice, the consequences of our alternative options are often collectively no better or worse than each other. In the second, I consider the potential implications of this conclusion. I propose that we should give predictable consequences less weight relative to other prudential and moral considerations and that we should feel less regret about the visible harms that we know we have caused.

The Case for Parity

Parity Conjecture: In cases where we deliberate between two or more alternatives, the total consequences of those alternatives are often no better or worse than each other.

My case for the Parity Conjecture relies on four assumptions.

First, our actions have widespread and varied consequences that largely occur in the unpredictable future. The consequences of our actions have their own consequences, which in turn have their own consequences. If we had acted differently, most downstream events would likely have happened differently.

Even very modest changes to our daily routine can cause further changes, and modest changes occasionally result in more significant effects. The choice of cereal for breakfast may seem an insignificant act, but it may influence what you choose to have for lunch, which may affect where you park your car, which may affect where dozens of other people's parking spots. This may in turn affect whose car is broken into, which may affect where people choose to live, which may affect who they meet and who they marry, and so on. An action may cause multiple effects, and each effect may cause multiple further effects, branching further outward over time. So long as the differences produced by alternative

actions grow slightly more significant or slightly more widespread over time, they are likely to eventually affect nearly every corner of human history.

Lenman (2000) builds a case that the consequences of many actions grow in perpetuity. His argument leverages the arbitrariness of parenthood and its effect on identity. Given that we have our parents essentially, any action that changes who procreates with whom (or when) will likely affect a whole ancestral chain. Through the growing collection of descendants, it will affect the world overall. It is relatively easy to indirectly affect identities in this way, so many of our actions have numerous and significant consequences.

Second, events can be intrinsically valuable in many different ways. It is good that people are happy and that they flourish in their intellectual pursuits. It is also good that people are treated with justice and fairness, that natural environments are protected and diverse species thrive, and that people are able to lead different sorts of lives, to follow different traditions, and to produce different cultural artifacts. They can come apart. An action's unpredictable future effects may include a boost to the prevalence of intellectual pursuits at the cost of crass pleasures. Or it may increase human welfare at the cost of biodiversity.

These distinct categories of value cannot be reduced to a single unified kind. The different sources of value cannot be construed as promoting the same basic end in different ways. Values are too diverse. The difficulties we face in evaluating tradeoffs between different sources of value precisely attest to that. If value ultimately reduced to a single fundamental kind, we would need only to know the reductive facts to make comparisons trivially easy. However, nothing we could learn could make precisely comparing different sources of value trivially easy.

Even those who accept one fundamental kind of value should acknowledge this difficulty. Consider the many forms of pleasure: there are crass pleasures and intellectual pleasures, intense brief pleasures and continuous mild pleasures, pleasures of each different sensory modality. Even pleasures of one modality that take very different forms. Suppose that all that matters is how good things feel; can we precisely compare the pleasures of a cold drink on a hot day, of a hot cup of strong coffee, of a fancy chocolate cordial, and of a comforting bowl of soup along one fundamental dimension? Or can we precisely compare them with the auditory pleasures of a catchy melody, a meditative chant, or an energizing beat? The same challenge that suggests a diversity of basic values – the ambivalence we face in making precise distinctions – affects each superficially unitary value. Even if pluralists are wrong about the fundamental categories of value, the prospects of strict unitarianism are dim.

Third, comparisons between different categories of values are often rough. This means that alternatives that excel with respect to different intrinsic values are often on a par in Ruth Chang's (2002) sense of the term. Two options are *on a par* if they are inexactly equivalent: neither is better or worse than the other but a relatively modest improvement to one option won't tip the balance in its favor. Alternatives are often on a par because the tradeoffs between sorts of value are imprecise. What exact number of headaches is William Shakespeare's corpus worth? What future human population reduction would be justified to preserve the manatee? The challenge of these questions is not just epistemic. There is nothing to learn about Shakespeare, headaches, or manatees that would make them straightforward. In principle, these questions appear to admit of no precise answers. The difficulty we face in assessing certain tradeoffs reflects some fuzziness of the values themselves.

If two possible sets of widespread consequences are on a par, then not only is

neither is more valuable than the other, but a minor tragedy here or there added to one set probably won't make it less valuable than the alternative. If a detailed accounting of each of the possible future courses of history that start from our choices won't tell us which choice is better or worse than another, adding in an additional murder (let alone a lost wallet) to one history is unlikely to settle the matter.

Fourth, viewed abstractly, the futures produced by our alternative actions are likely to differ unsystematically and mostly in a large number of relatively subtle ways. The futures may be very different in their details but closely resemble each other in terms of broad statistics. One future may have a few more murders but slightly greater artistic achievements (while of course having totally different victims and artworks). One future might preserve more biodiversity but have a slightly lower population.

Alternative choices are likely to excel with respect to different values once all of their consequences are considered. Few actions with widespread consequences are better than their alternatives in all ways. This assumption is motivated by the thought that in the grand scheme of things, the course of history is directed mostly by very general trends. An individual's actions may introduce small alterations to the future that affect every detail, but they are unlikely to shift the overall trajectory. If they don't shift the overall trajectory, many of their effects will be better or worse through happenstance, and it is unlikely they will all fall in the same direction.

Consider the effect of adopting a dog from a shelter. Without the adoption, the adopted dog would have ultimately gone to a different family. That family may instead have adopted a different creature. The chain of switched adoptions may extend far into the future. Some dogs would be better matched in some ways and some worse in others. Some would have better fits of personality. Some

who will die young in an accident would live to old age. Still, it is implausible that your adoption would by happenstance lead to systematically better lives for hundreds of other dogs.

This final assumption is particularly plausible if we think of our actions not as altering the course of history from one definite path to another, but as shifting the probabilities of a large family of possible futures. Our actions may foreclose on radically different possible futures without altering the general distribution of remaining possibilities. If these large families are composed of specific possibilities that deviate arbitrarily from a shared trend, a fair assumption, then the consequences of any two actions are highly likely to resemble each other. The sets of possible futures can be thought of as akin to large random samples from a single population: such samples are likely to closely resemble each other and differ only in unsystematic ways.

A deeper defense of these assumptions is beyond the scope of this paper. Because each is a reasonable conjecture, it is worth considering what implications they jointly have.

The argument for the Parity Conjecture is assembled as follows. Our actions are likely to indirectly contribute to widespread changes to future world history but are not likely to radically alter its general trend. Any action that affects most of the future without altering general historical trends is likely to have consequences that are better in many ways and worse in roughly as many others. These consequences mean that the future(s) produced by an action will excel with respect to some kinds of value and lag with respect to others, but probably not by very much. If categories of values are only roughly comparable, the futures produced by our actions will often be on a par.

Assuming that the unpredictable effects produced by our actions are widespread

and on a par, the total effects are likely to be as well. The predictable consequences of our actions will generally be relatively minor compared with their total consequences (Elgin 2015). Given the resilience of parity to slight modifications of the options, this means that the total consequences of our actions would also likely be on a par.

Implications of The Parity Conjecture

In light of the role that the consequences of our actions typically have in moral deliberation, acceptance of the Parity Conjecture should push us to rethink our moral responsibilities. I will now explore three potential implications of the Parity Conjecture: regarding tradeoffs with prudential considerations, regarding tradeoffs with other moral considerations, and finally regarding how much remorse it is appropriate to feel about the visible results of our actions.

Trade-offs with Prudential Considerations

We must balance our obligations to others with the pursuit of our own interests. How much sacrifice we owe it to others to make of own interests depends at least partly on what we can do for them. Accepting the Parity Conjecture may justify us in paying more attention to our own interests.

Lenman (2000) suggests a similar conclusion in response to our cluelessness about the consequences of our actions. According to Lenman, our cluelessness is such that each action we might take will have about a 50-50 probability of making things better or worse overall, no matter its predictable consequences. Given our inability to reliably make things better or worse, Lenman wonders, shouldn't it be permissible to just forget about moral demands and pursue our own projects?

Such radical unpredictability is consistent with the expected value of some actions being higher than others. Even if it is a toss-up whether our actions will make things better or worse, the expected value of our actions may still be high or low, and may correspond nearly perfectly with their predictable consequences. Consider: you can either bet your life savings on heads in a flip of a coin, or you can be handed \$100 and then bet your life savings on tails. While both options are likely to have you come out richer or poorer half the time, choosing between the two, you would be wise to take the \$100. While it may be a toss-up whether our actions will make the world better or worse, we can still be moved by the fact that the expected value is higher.

The idea that the Parity Conjecture supports favoring prudential considerations is more compelling.

Consider Peter Singer's classic case in which we are asked whether we should save a child from drowning in a shallow pond at the cost of the expensive clothes we happen to be wearing (Singer 1972). This case is widely regarded as strongly suggesting that morality is very demanding: given our ability to use our wealth to improve the lives of the impoverished, living a rather ordinary life in the developed world is an act of unacceptable negligence (Unger 1996).

Suppose that saving the child from drowning would not only come at a great cost, but would lead to a future in which the slight reduction in children drowning was balanced by a modest rise in spousal abuse. Suppose that these different harms and benefits are on a par. It is less plausible that we would still have a strong moral obligation to ruin our clothes to save the child.

Parity isn't guaranteed; if we see that there is only a small probability that the consequences of our alternatives will be on a par, it might make sense to downgrade our moral obligations proportionately. So if the prospect of parity

doesn't erase the fact that we have moral obligations, it may justify us in giving some additional preference to our own projects. In the shallow pond case, imagine that you know that if you do save the child, there is a 99.8% probability that the reduction in drownings will be balanced out by a rise in spousal abuse. Doesn't that make it at least a little more reasonable to favor your clothes? If it is true for 99.8%, it should also be true, to a lesser extent, for a 26% probability of parity. There is no clear cutoff in which the probability of parity should stop mattering.

These are plausible implications of the Parity Conjecture, but we shouldn't accept them too hastily. There are a few complications.

The first complication is that the significance of parities on decision making is largely unexplored and it isn't clear how we should balance the prospect of parity with the value we see in predictable tradeoffs.

Caspar Hare (2010) presents a puzzle regarding the way that predictable values should influence our prudential decisions when we know that total values are on a par. His discussion revolves around a simple thought experiment in which we are asked to choose between two mystery boxes. We do not know precisely what is in each box but we do know that the contents are on a par. The value is too different to strictly compare and if we were faced with the choice between them we would be ambivalent about which to choose. One of the mystery boxes also includes a small inducement. We know what this inducement is, and it is something that we value (e.g. \$1), but we also know that it is too small a thing to overcome the parity between the uncertain contents. The inducement introduces an asymmetry between the boxes. In light of the parity of the contents, it is tempting to let it break the tie.

Hare describes two decision theories that offer conflicting verdicts. One suggests

we should focus on what reasons we have and bracket the unknown. If the boxes are equal mysteries, the inducement is a tiebreaker. The other suggests that we should defer to the judgements we know we would make in light of full information. If we could open the boxes and examine the contents, we would be ambivalent about our choice given the overall parity. Since we know this, we should be ambivalent before opening the boxes even in our ignorance.

It is not clear to Hare or me which theory is right. Both have unintuitive consequences (e.g. Doody 2019a, 2019b).

Hare's theories obviously bear on the implications the Parity Conjecture, but they may not be directly applicable to the moral domain, or to cases where we must weigh tradeoffs between the moral and prudential domains. Prudential decision making may follow different rules from moral decision making. So it may be that we should be swayed by the inducement in Hare's cases but not be swayed by a corresponding inducement in cases of moral parity (or vice versa). Or it may be that we are licensed to ignore moral inducements between options on a moral par in favor of satisfying our interests, even when morality alone favors the inducement.

Still, the theories suggest two responses to the Parity Conjecture. On the first, accepting the Parity Conjecture shouldn't change how we think about how the predictable consequences figure into moral deliberation. We should be swayed by the predictable consequences because they introduce an asymmetry of our reasons and they are all we have to go on. On the second, accepting the Parity Conjecture between the overall consequences would rob the relatively small predictable consequences of their significance. On this view, the predictable consequences mislead. We know that they are balanced out by unpredictable consequences so that the total consequences are likely on a par.

The second complication is that it is not obvious how to think about the significance of moving from an option that is on a par to an option that is slightly better. If one action we might take makes things better than another, but only marginally beyond being on par, how should we think of it? This is a tricky question, in part because parity seems rarely to admit of obvious borderline cases. It will be hard to find cases that are clearly on a par where a relatively small inducement will make one option clearly better, so hard to assess our intuitions about the normative difference.

While it is possible that the idea that one set of consequences could be minimally better than another is confused, parity appears to be compatible with a threshold of minimal betterness existing and reflecting a major normative difference. In that case, even if our options are likely to be on a par, it is possible that they won't be; the chance we have to make a big difference may dominate our reasons for acting one way or the other.

Perhaps we should think of our dilemma as akin to that of the consumer whose morally questionable purchases are unlikely to influence any company's choice to maintain or expand its repellent practices. Such decisions need to take into account the small chance of making a major difference by tipping things past a threshold (Kagan 2011). This small chance of a significant effect may be enough to decide the matter even if the most likely outcome is that our action will not make any difference. If we should think of the chance of crossing the threshold from parity between alternatives to a minimal non-paritous difference as being like the consumer's choice, then our moral reasons might not be much affected by the probability of parity.

While this is possible, it seems unlikely. If two options are on a par, then a slight improvement to one may suffice to make it slightly better, but it shouldn't make it much better than its alternative. Just as parity suggests rough equivalence, a

small improvement over parity should suggest a very small improvement over rough equivalence. Such changes shouldn't be treated as significantly different from a small definite improvement over an exactly comparable option.

Trade-offs with Other Moral Considerations

The Parity Conjecture may have significant implications for weighing conflicting moral considerations. If producing visibly good consequences is unlikely to correlate with making things better overall, we should perhaps grant predictable consequences less weight when trying to decide what we should do.

Suppose that the value of consequences competes with deontic considerations in moral deliberation. We have some responsibility for preventing bad consequences and encouraging good consequences, but we also have reason to treat others as ends rather than means, to act in ways that we could reflectively endorse for everyone, and to manifest virtues of character. The fact that the consequences of our actions are often on a par means that they shouldn't play the role they typically play in moral deliberation. But since this leaves the other considerations unaffected, they should come to have a greater relative bearing on how we act.

Perhaps keeping our word matters at least a little, whatever the consequences. Sometimes keeping our word requires significant and predictable sacrifices. Those sacrifices in isolation would make it clearly wrong to keep our word. But the fact that those predictable sacrifices are dwarfed by unpredictable consequences that are very likely to be on a par may mean that we should keep our word even when we see that it will do great immediate harm.

Suppose that a person promises to take their friend to a gun show before hearing about that friend's recent aborted suicide attempt. The gun show might provide means or inspiration for their next attempt – should the promise be broken?

Aiding in some way in the friend's death would be a bad thing by itself, but it would also have causal ripples into the far future. Suppose that keeping and breaking the promise each have consequences that are ultimately on a par. If the duty to keep our word counts for something, and isn't merely one more good consequence to add to the lot, then it may tip the balance of reasons in favor of keeping the promise.

Alternatively, perhaps the promiser shouldn't keep their word, but the reasons are different than consequentialists maintain. The problem isn't that the action is harmful. Most actions are harmful in some ways and not in others, and, if the Parity Conjecture is true, on the balance the results of alternative actions are often on a par. The problem with keeping the promise is instead that it harms *a friend*, and though our special moral obligations to our friends don't trump the obligations we have to promote utility overall, they are sufficient to often act as a tiebreaker in cases of likely parity.

If parity robs predictable consequences of their significance, it may mean that our moral reasons are dominated by non-consequentialist considerations. This wouldn't be because consequentialists are completely wrong about the nature of normativity. They may be right that insofar as consequences do differ significantly between our actions, we ought to be swayed by the greater consequences. But the causal interconnectedness of reality may undercut the significance of our intuitions about the morality actions given their predictable consequences in isolation.

Appropriateness of Remorse

Finally, accepting the Parity Conjecture may change how it is we should feel in retrospect about the visible good and bad consequences of our actions. When our actions have obvious harmful consequences, it is generally appropriate to feel

guilt, regret, or remorse about what was done. These feelings are particularly apt when the bad consequences are not balanced or outweighed by other good consequences. Though it is still often appropriate to feel bad about the harms one has caused in the pursuit of the greater good, finding out that the visible harms of our actions are balanced by equally weighty benefits, and so the consequences are not net bad overall, is something of a relief.

If our actions have widespread unpredictable consequences, then those consequences are likely to include both good and bad effects. If the Parity Conjecture is true, these negative and positive consequences will very often balance each other so that the total consequences of the actions that we took are no worse than the actions we could have taken. Even if we know that some of the consequences of our actions are bad, we can rest assured that it is likely that the unknown consequences of our actions make up for them.

We can console ourselves with the thought that it is likely that if we could see all of the effects of our actions, we would be ambivalent about whether we made things better or worse. Accepting the Parity Conjecture may not completely relieve us of our guilt about the harms we cause, but it should somewhat dampen the regret we feel.

Conclusion

The probability and significance of the Parity Conjecture both remain unclear, but the Parity Conjecture is plausible and the potential implications are great.

Appendix 1: Modeling the Probability of Parity

How should we evaluate the probability that the consequences of two actions are on a par? This paper offered an abstract argument that parity might be surprisingly common. That argument left it open precisely how common that is. This appendix aims to clarify this question with the help of a model.

One way to approach questions as empirically uncertain as the Parity Conjecture is to attempt to model the uncertain phenomena. If we can develop a framework for building models, we can separate out parameters for those aspects of the theory we are uncertain about and we can see how different values for the parameter affect the results of the models.

In order to investigate parity with models, we need some way of abstractly representing the effects actions have on different kinds of value and the way that tradeoffs between kinds of value can fail to be determinate.

For this purpose, we can concentrate on modelling the overall effects an action has within each category of value. For each category, we can assign the action a number that represents the bearing the consequences of the action have on value of that kind. Higher numbers represent better consequences, but the imprecision of tradeoffs prevents us from drawing straightforward conclusions about the relative value of different categories.

Let's call the array of numbers representing each action's effect on a value category the model's 'value score'.

A value score for an action might be

[32, 50, 21, 99, 6]

and represent

32 - The value of the action's effect on biodiversity

- 50 - The value of the action's effect on human welfare
- 21 - The value of the action's effect on human achievements
- 99 - The value of the action's effect on justice
- 6 - The value of the action's effect on human knowledge

The numbers comprising a value score indicate something about the relative significance of the effect on different values. To allow for the imprecision of value tradeoffs, we also include in each model a range (the "comparison range") representing the extent to which values of different categories are on a par in terms of axiological significance. The wider the range, the greater the fuzziness, and the greater that the difference in the numbers assigned to distinct categories in the value score must be in order for an effect on one value category to be more significant than an effect on another value category.

A comparison range might be

[.9, 1.1]

and signify that the numbers representing each category of value are roughly as important as the numbers representing each value of another category that is about 20% greater or lesser. (The 20% allowance results from the fact that one value might be 10% greater than its precise score indicates while the other is 10% less.)

With a value score and a comparison range in place, we can define parity in terms of a weighting scheme. A weighting scheme includes a number in the comparison range for each category of value and represents the values' relative significance according to that precisification.

We can think of weighting schemes as permissible precisifications of the relative value of value categories. A weighting scheme might have biodiversity in the upper range of its plausible precisifications by weighing it at 1.1, while it has the value of justice is at the lower range of its plausible precisifications by weighing it at .9.

Two models are on a par when there exists both a weighting scheme in which the first comes out as less and another weighing scheme in which it comes out as more. In contrast, one model is strictly better than another when it comes out as better for every weighting scheme.

Among models with two categories of value and weighting scheme [.9, 1.1], the model with value score

$$[50, 0]$$

is better than the model with value score

$$[0, 40]$$

because

$$50 \times \alpha + 0 \times \beta$$

is greater than

$$0 \times \alpha + 40 \times \beta$$

for any α, β between .9 and 1.1

but among models with comparison range $[.8, 1.2]$, the model with the value score

$$[50, 0]$$

is on a par with the model with the value score

$$[0, 40]$$

because

$$50 \times \alpha + \\ 0 \times \beta$$

is greater than, and less than

$$0 \times \alpha + \\ 40 \times \beta$$

for $\alpha, \beta = 1$ and $\alpha = .8, \beta = 1.2$ respectively

Suppose that actions have arbitrary effects on all of the categories of value in the future and that those effects are independent of each other. We can then measure the probability of parity by considering how likely it is that randomly generated models are on a par. This is easy to do with a computer program, and the results generated below were produced with such a program (available [redacted]).

For these results, I chose several selections of parameters (the numbers of value categories and comparison ranges) and repeatedly generated pairs of models using random numbers to fill in each model's value score. Each pair of generated

models was checked for parity. Repeated enough times, this procedure can establish the probability of parity given the assumptions of the model with a great deal of precision.

The probability of parity between two randomly generated models increases both with the size of the comparison range and the number of value categories. Increasing the size of the range allows for greater distortions of the generated values in each model, which makes it easier for one model to surpass the other under some weighting of values.

We can find a weighting scheme that favors one model over another by increasing the weight of the values in which it leads the other and decreasing the weight of the values in which it lags. The more value categories, the more likely that each model leads the other in a number of dimensions. Additionally, increasing the number of different value categories smooths out the chanciness of the random assignments, making it less likely that a probabilistic aberration strongly favors one model over another.

The results of following this procedure show that for small ranges and small numbers of distinct value categories, two random models are likely not to be on a par. However, parity is a significant possibility for pairs of models with more than one value category, even when the comparison range is fairly narrow. For models with more than a hundred value categories, parity is the norm. For models with more than a thousand value categories, parity is the rule, even for very narrow comparison ranges.

Table 1: The probability of parity by comparison ranges and number of value categories rounded to the nearest percentile.

	.9 - 1.1	.8 - 1.2	.7 - 1.3	.6 - 1.4	.5 - 1.5
3	9%	18%	27%	35%	43%
10	19%	36%	53%	66%	77%
100	57%	88%	98%	100%	100%
1000	99%	100%	100%	100%	100%

If we think that such models and their generation procedure approximately reflect the situation we face in choosing between actions, we should take the possibility of parity in outcomes very seriously.

Bibliography

- Chang, Ruth. 2002. “The Possibility of Parity.” *Ethics* 112 (4): 659–88.
- Doody, Ryan. 2019a. “Opaque Sweetening and Transitivity.” *Australasian Journal of Philosophy* 97 (3): 559–71.
- . 2019b. “Parity, Prospects, and Predominance.” *Philosophical Studies* 176 (4): 1077–95.
- Elgin, Samuel. 2015. “The Unreliability of Foreseeable Consequences: A Return to the Epistemic Objection.” *Ethical Theory and Moral Practice* 18 (4): 759–66.

- Hare, Caspar. 2010. "Take the Sugar." *Analysis* 70 (2): 237–47.
- Kagan, Shelly. 2011. "Do I Make a Difference?" *Philosophy & Public Affairs* 39 (2): 105–41.
- Lenman, James. 2000. "Consequentialism and Cluelessness." *Philosophy & Public Affairs* 29 (4): 342–70.
- Singer, Peter. 1972. "Famine, Affluence, and Morality." *Philosophy & Public Affairs*, 229–43.
- Unger, Peter K. 1996. *Living High and Letting Die: Our Illusion of Innocence*. Oxford University Press, USA.